

CHAPTER 1

INTRODUCTION

Urbanization is one of the most significant global trends of the 21st century, fundamentally reshaping landscapes, economies, and human lifestyles. With rapid population growth and urban sprawl, cities are expanding at an unprecedented rate. This transformation poses complex challenges for urban planners, policymakers, and environmentalists who must balance development with sustainability. Accurate and timely analysis of urbanization patterns is essential for informed decision-making, efficient infrastructure development, and mitigating negative impacts such as congestion, pollution, and loss of green spaces.

Traditionally, urbanization analysis has relied on census data, manual mapping, and remote sensing interpretation, which are often time-consuming and lack the granularity or timeliness needed for proactive planning. In recent years, advances in artificial intelligence—particularly deep learning—have enabled powerful tools for automated image analysis. Among these, convolutional neural networks (CNNs) have shown great success in image segmentation and classification tasks. One such architecture, U-Net, has emerged as a leading model for semantic segmentation, especially in the context of biomedical and satellite imagery.

This project introduces the Urbanization Chat Bot, a system that leverages U-Net for semantic segmentation of satellite imagery to predict urban expansion patterns. The core functionality includes preprocessing satellite images, training a U-Net model to identify urbanized areas, and visualizing the predicted transformation over time. Additionally, a chatbot interface is integrated into the system to make it interactive and user-friendly, enabling users to ask questions and receive predictions or visual feedback in a conversational format.

By combining image segmentation techniques with natural language processing, the Urbanization Chat Bot aims to democratize access to urban data insights, offering both technical and non-technical users a platform to explore and understand urban growth trends. This project not only demonstrates the practical application of deep learning in geospatial analysis but also sets the foundation for future enhancements, such as real-time augmented reality interfaces or integration with GIS systems.

1.1 AIM

The primary aim of this project is to design and implement an intelligent and interactive system, named **Urbanization Chat Bot**, that leverages deep learning techniques to analyze and predict patterns of urban growth using satellite imagery. The system seeks to bridge the gap between complex geospatial data and user accessibility by combining a convolutional neural network-based segmentation model (U-Net) with a conversational AI interface. By doing so, the project aims to enable users—from urban planners and researchers to students and policymakers—to intuitively interact with the model, retrieve insights, and visualize the transformation of land over time due to urban expansion.

This project addresses the growing need for automated, scalable, and user-friendly tools capable of understanding spatial changes in urban environments. It aims not only to provide accurate predictions through satellite image segmentation but also to present those results in a clear, interactive format that supports informed decision-making. Additionally, the chatbot interface is designed to democratize access to these predictive insights, allowing users to query the system in natural language without requiring technical expertise in machine learning or geospatial analysis.

1.2 MOTIVATION

Urbanization is a transformative force that is reshaping cities and societies around the globe. As more people migrate from rural to urban areas in search of better opportunities, infrastructure, and lifestyles, the demand for effective urban planning and sustainable development becomes increasingly urgent. However, monitoring and predicting urban expansion is a complex task that requires continuous observation, analysis of spatial data, and the ability to forecast future growth accurately. Traditional methods such as manual mapping, periodic surveys, and static GIS systems often fall short in terms of efficiency, scalability, and real-time responsiveness.

The motivation for this project arises from the need to modernize and enhance the way we understand urbanization. With the availability of high-resolution satellite imagery and advancements in machine learning, there is an opportunity to automate the detection and analysis of urban growth patterns with high accuracy and speed. Deep learning models, particularly convolutional neural networks like U-Net, have proven to be highly effective in semantic segmentation tasks—making them ideal for identifying built-up areas from satellite data.

In parallel, the increasing use of conversational interfaces in applications across industries highlights the value of integrating AI-driven insights into user-friendly platforms. By embedding a chatbot that allows users to ask questions, request visualizations, and interpret results, this project not only improves accessibility but also broadens the reach of urban analysis tools to non-technical audiences. This is particularly beneficial for city planners, researchers, students, and even citizens who want to understand how urbanization is impacting their environment.

Furthermore, real-time decision-making in urban development often requires fast and interactive tools. The fusion of deep learning for prediction and chat-based interfaces for communication forms a novel solution that meets the growing demands of smart cities and data-driven governance.

1.3 PROBLEM STATEMENT

Urbanization is an ongoing and accelerating global phenomenon with significant implications for infrastructure development, environmental sustainability, and socio-economic planning. As cities expand and new urban areas emerge, it becomes increasingly critical to monitor and analyze patterns of land use change. However, accurately predicting and understanding urban transformation remains a major challenge due to the complexity, scale, and dynamic nature of urban growth.

Traditional approaches to urban analysis rely heavily on manual interpretation of satellite imagery, periodic surveys, and geographic information systems (GIS), which are often labor-intensive, time-consuming, and unable to provide timely updates. Moreover, these methods usually require specialized expertise, making it difficult for non-technical stakeholders such as local authorities, students, or citizens to engage with urban planning data in a meaningful way. The absence of an automated, scalable, and user-friendly solution significantly hampers real-time decision-making and limits the potential for data-driven urban planning.

Additionally, while machine learning techniques—particularly deep learning—have shown promising results in tasks like image segmentation, their deployment in real-world applications is often constrained by usability barriers. Most deep learning tools for satellite image analysis are designed for researchers or developers, with minimal focus on accessibility or interactivity. There exists a gap in integrating these powerful models into systems that are intuitive and interactive for general users.

Furthermore, existing urban prediction tools typically lack a mechanism for real-time user engagement. They often present static outputs with no interactive querying capability, which limits the flexibility and responsiveness needed in real-world applications. As a result, users are unable to ask questions, explore alternative scenarios, or extract tailored insights from the data, which restricts the usefulness of such tools in planning, education, or public awareness.

This project addresses these critical gaps by proposing an intelligent, conversational system—**Urbanization Chat Bot**—that combines a deep learning-based semantic segmentation model (U-Net) with an interactive chatbot interface. The goal is to provide users with an accessible platform to input satellite images, receive predictive analysis of urban growth, and engage with the system using natural language. By doing so, the project aims to make urban analysis not only more accurate and automated but also more approachable and user-centric.

1.4 EXISTING SYSTEM

In the field of urban growth prediction and land-use monitoring, several systems and tools have been developed over the years to analyze satellite imagery and generate insights for planning and environmental management. Traditional systems often rely on Geographic Information Systems (GIS), manual annotation, and classical image processing techniques to identify changes in land cover. While these tools have been instrumental in historical analysis, they come with significant limitations in terms of scalability, automation, and real-time usability.

One of the widely used approaches involves supervised classification of satellite imagery using pixel-based or object-based techniques. Tools like ERDAS Imagine, ENVI, and QGIS support such workflows where users must manually define training samples and apply classifiers like maximum likelihood or support vector machines (SVM). However, these methods are highly dependent on expert knowledge, prone to human error, and can be time-consuming when processing large datasets.

In recent years, advancements in remote sensing and artificial intelligence have given rise to more sophisticated systems powered by machine learning and deep learning algorithms. Convolutional Neural Networks (CNNs), especially architectures like U-Net, SegNet, and DeepLab, have demonstrated exceptional accuracy in semantic segmentation tasks. Some modern platforms, including Google Earth Engine and IBM PAIRS, offer cloud-based satellite image analysis using such models. However, many of these platforms are either not open-source, lack interactive capabilities, or require significant technical expertise for operation and customization.

Despite these developments, existing systems often fall short in two major areas: accessibility and interactivity. Most deep learning-based solutions are built for research purposes, offering limited user interaction and visualization tools. They typically do not include conversational interfaces or support for non-technical users to engage with

predictive models in a meaningful way. Furthermore, outputs are often presented as static images or reports, without mechanisms for dynamic querying or guided analysis.

There is also a lack of systems that combine predictive modeling with user-centered design, especially chat-based interfaces that allow real-time interaction. This presents a significant opportunity to enhance user experience by embedding deep learning predictions within a natural language processing (NLP) framework. Such integration could make complex data analysis tools usable by a wider audience, including policymakers, educators, and the general public.

Therefore, while existing systems provide valuable capabilities for satellite image processing and urbanization analysis, they do not adequately address the need for a fully interactive, AI-driven, and accessible solution. This gap highlights the importance of developing the proposed Urbanization Chat Bot, which aims to bridge this divide by integrating deep learning-based image segmentation with a conversational interface to deliver urban transformation insights in a user-friendly and interactive manner.

1.5 PROPOSED SYSTEM

The proposed system, titled **Urban Shift Dynamics**, is designed to provide an intelligent, interactive platform for analyzing and predicting urban transformation using satellite imagery. Unlike traditional static analysis tools, this system integrates deep learning with a conversational AI interface to make urban growth insights accessible to both technical and non-technical users. The core of the system is built on the U-Net architecture, a convolutional neural network specifically tailored for semantic segmentation tasks. This enables the model to detect and classify urbanized regions within satellite images with high accuracy and spatial precision.

The system workflow begins with the ingestion of satellite imagery data, which is preprocessed for noise reduction, resizing, and normalization to ensure consistency and

model compatibility. The processed images are then fed into the U-Net model, which segments the input into urban and non-urban classes. The output is a mask that visually represents predicted urban areas, which can be overlaid on the original image to show transformation patterns over time.

To enhance usability, a chatbot interface is integrated into the system, allowing users to interact with the model using natural language queries. For example, users can ask, “Show me how this area has urbanized over the last decade” or “Predict urban expansion for this region.” The chatbot processes these queries, communicates with the backend model, and responds with relevant predictions, explanations, or visual outputs. This natural language interface significantly lowers the barrier for engaging with the system, making it suitable for urban planners, researchers, educators, and the general public.

The system is implemented using TensorFlow for model development, OpenCV for image handling, and Python-based frameworks for backend integration. The chatbot functionality is developed using natural language processing (NLP) tools such as Rasa or similar platforms to interpret user input and manage dialog flow. A modular architecture ensures that the components are easily maintainable and scalable, allowing for future enhancements such as real-time prediction, multilingual support, or integration with GIS databases.

By combining the predictive power of deep learning with the simplicity and intuitiveness of a chat-based interface, the proposed system offers a novel solution to the challenges faced by existing urban analysis tools. It aims to democratize access to urban transformation insights, support smarter city planning, and promote data-driven decision-making in a rapidly urbanizing world.

CHAPTER 2

LITERATURE SURVEY

[1] **Dilraj Bhinder and Dr. Devaraj Verma (2023) , "Literature Review on Change Detection Using Remote Sensing Imagery"**, This paper presents an extensive review of change detection (CD) techniques in remote sensing imagery, examining both traditional and advanced approaches. The authors outline the key steps in CD processes, such as image registration, feature extraction (color, texture, shape), and various analysis methods using statistical and machine learning techniques. A significant contribution of this review is its discussion on the use of Support Vector Machines (SVM) for classification. Moreover, it addresses existing challenges like limited training datasets and dependence on manual feature engineering. The paper emphasizes the role of modern solutions such as transfer learning and Siamese networks in improving classification accuracy and automation. This work serves as a foundational understanding of the landscape of CD techniques and their evolution.

[2] **Sebastian Hafner, Heng Fang, Hossein Azizpour, and Yifang Ban (2024), "Continuous Urban Change Detection from Satellite Image Time Series with Temporal Feature Refinement and Multi-Task Integration"**, This paper introduces a novel approach for continuous urban change detection using satellite image time series. The authors leverage a deep learning model based on **Temporal Feature Refinement (TFR)** and **Multi-Task Integration (MTI)**. The TFR module models temporal dependencies in satellite image sequences, while the MTI module uses Markov networks to jointly perform classification and segmentation. This integrated model helps reduce the overhead of using multiple separate systems. One of the notable advantages is its ability to perform real-time monitoring of urban development. However, the increased model complexity demands more computational resources and sophisticated data preprocessing. The paper stands out for its holistic approach to modeling temporal urban change, which has applications in urban planning and policy-making.

[3] Saurabh Srivastava and Dr. Tasneem Ahmed (2023), "An Approach to Monitor Urban Growth through Deep Learning based Change Detection Technique using Sentinel-2 Satellite Images", This study utilizes Sentinel-2 satellite imagery from 2017 to 2022 to monitor urban growth through a deep learning-based change detection framework. The SNAP tool is used for preprocessing the images, including atmospheric and radiometric corrections. The authors implement classifiers such as ANN, SVM, and MD to differentiate between urban and non-urban areas using indices like NDBI, NDVI, and BUI. The method enables accurate urban classification and efficient monitoring of urban expansion. The integration of tasks reduces computational costs compared to running separate models. The research demonstrates the practical applicability of supervised classification techniques and emphasizes the significance of remote sensing tools in understanding urban dynamics.

[4] Xiaoxue Feng, Peijun Li, and Tao Cheng (2021), "Detection of Urban Built-up Area Change from Sentinel-2 Images Using Multi-Band Temporal Texture and One-Class Random Forest", This paper focuses on detecting urban changes using Sentinel-2 imagery from 2015 to 2019. The authors derive multi-band temporal textures through pseudo cross multivariate variograms and combine these with spectral features and NDVI indices. An improved One-Class Random Forest (IocRF) algorithm is used to classify urban changes. The study highlights the effectiveness of temporal feature refinement in minimizing noise in satellite images, making the analysis more reliable. However, the method's limitation lies in its scalability—it may not be ideal for large-scale or global datasets. This paper contributes to the development of texture-based classification techniques and provides a focused methodology for urban analysis at smaller scales.

[5] Mandicou BA, Pape Ibrahima Thiam, Etienne Delay, Charles Abdoulaye Ngom, Idy Diop, and Alassane Bah (2024), "Deep Learning-based Land Use and Land Cover Changes Detection from Satellite Imagery: A Case Study of the City of Richard Toll", The study explores land use and land cover (LULC) changes using satellite images from Google Earth Pro, segmented into eight land classes. Change detection is achieved by comparing segmentation masks across different time periods. While the method excels in identifying subtle land use changes, it struggles with handling continuous or real-time data streams. This research demonstrates how deep learning can aid in environmental monitoring, especially in smaller urban regions. Despite its limitations in scalability and real-time adaptation, the approach successfully captures incremental urban transformations, which are essential for regional planning.

[6] Martin Leichter, Filip Sadlo, and Joachim Denzler (2023), "UrbanChange3D: A Spatiotemporal Dataset for Change Detection and Classification of Urban Outdoor Scenes", This paper introduces **UrbanChange3D**, a large-scale spatiotemporal dataset designed to support deep learning tasks in urban change detection. The dataset includes over 110,000 3D point cloud images from 11 major cities across four countries, with data collected between 2008 and 2023. The dataset offers per-pixel ground truth annotations for binary change detection, semantic segmentation, and semantic change detection. The authors address challenges in semantic variability over time and provide a benchmark for evaluating model performance. This resource is particularly valuable for training generalizable models and understanding long-term urban transformations. However, due to its massive scale, it requires significant computational resources and careful data handling. The dataset is expected to push the boundaries of research in urban scene understanding and automation.

[7] **Daisuke Sugai, Akira Oda, and Kazuhiro Fukui (2023), "Learning Spatiotemporal Features for Change Detection from Multispectral Images"**, In this paper, the authors propose a novel deep learning framework to detect changes in multispectral satellite images by learning both spatial and temporal features. Their model uses two encoders to process bi-temporal input images, which are then fused to extract meaningful temporal features. A decoder follows this step to generate a precise change mask. One unique contribution of the work is its **adaptive loss function**, which accounts for changes in spatial scale and noise. The approach is particularly useful in differentiating between genuine urban growth and changes due to environmental noise or sensor inconsistencies. Though promising, the method requires careful calibration and a diverse training dataset to ensure consistent performance across varied terrains and image resolutions.

[8] **Ghazal Roostaei, Julian Brabandere, Amirhossein Habibian, and Diederik Paul Moeys (2023), "Multiscale Change Detection Using Transformer and Convolutional Neural Networks"**, This paper presents a hybrid deep learning approach that integrates **Transformers** and **Convolutional Neural Networks (CNNs)** for multiscale change detection in satellite images. The authors address the limitations of purely convolutional methods in capturing long-range dependencies by incorporating Transformer blocks into the encoder. This enables the model to better learn both local textures and global context. The decoder then processes these features for accurate change segmentation. This architecture excels in detecting both small-scale and large-scale urban changes with high precision. However, the Transformer's high computational cost can be a drawback in low-resource settings. The method sets a new benchmark in fusing multiscale information, making it highly suitable for diverse urban monitoring tasks.

[9] **S. Maheswari, Dr. M. Kamarasan, and Dr. P. Subashini (2022), "Comparative Analysis of Deep Learning Techniques for Urban Change Detection Using Remote Sensing Images"**, This paper offers a comparative evaluation of various deep learning models for urban change detection, including CNNs, ResNet, and U-Net. The authors preprocess remote sensing images and apply these models to extract change features. The study finds that **U-Net** performs best in preserving spatial details and accurately detecting small changes. In contrast, ResNet is more efficient in processing time but may overlook fine-grained information. The authors discuss the trade-offs between model complexity, accuracy, and computational efficiency. The analysis helps guide practitioners in selecting appropriate models based on their specific goals and resource constraints. This work serves as a valuable reference for understanding the strengths and limitations of leading deep learning approaches in urban settings.

[10] **Lin Di, Zhenhong Jia, Yunchen Gao, Xiang Li, and Zhiqiang Deng (2023), "Cross-Temporal Domain Adaptation for Remote Sensing Image Change Detection"**, This study addresses the problem of domain shift in change detection by proposing a **Cross-Temporal Domain Adaptation (CTDA)** method. Remote sensing images often vary in appearance due to differences in sensors, lighting, or atmospheric conditions across time. The CTDA framework aligns feature distributions across temporal domains using adversarial training, improving model robustness. A novel loss function is introduced to balance classification and domain alignment tasks. The authors demonstrate the model's effectiveness on several benchmark datasets, showing significant improvements in generalization performance. The method is especially useful in real-world applications where ground truth labels for newer images are unavailable or limited. However, it demands additional training time and careful tuning.

CHAPTER 3

SYSTEM REQUIREMENTS AND SPECIFICATIONS

3.1 SYSTEM REQUIREMENTS

1. Hardware:

- **Processor:** Intel Core i5 (8th Gen) / AMD Ryzen 5 or equivalent
 - **RAM:** 8 GB
 - **Storage:** 256 GB SSD (or 500 GB HDD)
 - **Graphics:** Integrated GPU (sufficient for basic data visualization)
 - **Display:** 13” screen with 1366x768 resolution or higher
 - **Peripherals:** Keyboard, Mouse, Internet-enabled Device
- Recommended Requirements (Production / Hosting)**
- **Processor:** Intel Xeon / AMD Ryzen 7 or higher (for server)
 - **RAM:** 16–32 GB (for concurrent data processing and API requests)
 - **Storage:** 512 GB SSD or more (expandable based on dataset size)
 - **GPU (Optional):** NVIDIA GTX/RTX or equivalent (for deep learning-based image change detection, if used)
 - **Server:** Dedicated or cloud-based VM with auto-scaling support.

2. Software:

Frontend

- **Framework:** Angular 15 or later
- **Languages:** HTML5, CSS3, TypeScript
- **Dependencies:** Angular CLI, RxJS, Bootstrap / Tailwind CSS (if used)

Backend

- **Framework:** FastAPI (Python 3.9+)

Libraries:

- uvicorn (for ASGI server)

- pydantic (for data validation)
- scikit-learn, pandas, numpy (for data analysis)
- opencv-python or PIL (for image processing)
- reportlab / pdfkit (for PDF export)

Development Tools

- **IDE:** Visual Studio Code / PyCharm
- **Version Control:** Git with GitHub/GitLab
- **Package Manager:** pip, npm

Operating System

- **Development OS:** Windows 10/11, Ubuntu 20.04+, or macOS
- **Server OS:** Ubuntu Server 20.04 LTS or higher (for production)

3. Network:

Development

- **Internet Connection:** Minimum 5 Mbps (for package downloads and API testing)
- **Localhost testing:** <http://127.0.0.1:8000> (FastAPI) and <http://localhost:4200> (Angular)

Deployment

- **Hosting Platform:** AWS EC2, Azure App Service, or Heroku (for API and frontend hosting)
- **Bandwidth:** Minimum 10 Mbps per 100 concurrent users (scalable as needed)
- **Domain & SSL:** Custom domain with SSL certificate for secure HTTPS access
- **Firewall Configuration:** Allow inbound traffic on ports 80 (HTTP) and 443 (HTTPS)

3.2 SYSTEM SPECIFICATIONS

1. Functional:

- **Input:** User queries, filter selections, admin data uploads.
- **Processing:** Query interpretation, data analysis, PDF/report generation.
- **Output:** Chatbot responses, charts/graphs, downloadable reports.
- **UI:** Chat interface, dashboard, filters, and export options.

2. Non-Functional:

- **Performance:** Responses within 2 seconds; API latency under 500ms for standard queries.
- **Accuracy:** Minimum 85% model accuracy on urbanization queries; confidence-based handling.
- **Scalability:** Supports horizontal scaling; handles 10x traffic with containerized deployment.

3. User:

- **Ease of Use:** Simple and intuitive interface.

CHAPTER 4

SYSTEM DESIGN

4.1 SYSTEM ARCHITECTURE DIAGRAM

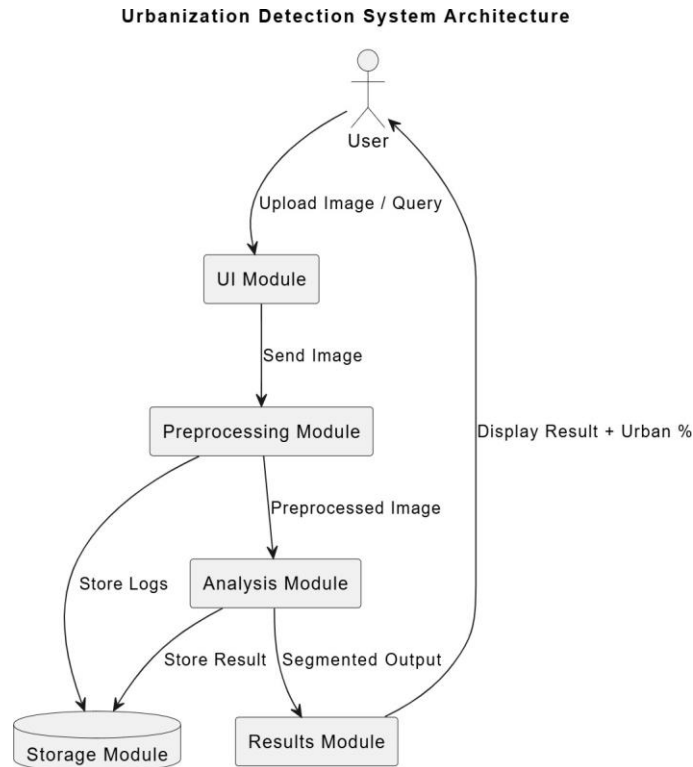


FIG 4.1 SYSTEM ARCHITECTURE DIAGRAM

The system architecture of *Urban Shift Dynamics* is designed as a modular and layered framework that enables smooth and intelligent interaction between users and the backend machine learning model. The architecture begins at the User Interface (UI), where users submit queries related to urbanization trends, migration patterns, or city growth metrics. These queries are sent to the FastAPI Backend, which functions as the central controller, routing requests to appropriate components for processing.

The backend passes the user input to the Machine Learning Model, which interprets natural language using pre-trained data and returns insightful responses. The model may also interact with a Database or Knowledge Base containing structured urban data, historical records, or analytical insights to enrich the chatbot's responses.

To support dynamic and accurate information retrieval, the system can connect with External APIs, such as those providing demographic, infrastructure, or environmental data. All system components, including the backend and ML model, are containerized using Docker for scalable deployment across different platforms and environments.

This architecture ensures high performance, accuracy, and scalability, enabling *Urban Shift Dynamics* to deliver timely, context-aware responses while adapting to increasing user demands and evolving urban datasets.

4.2 DATA FLOW DIAGRAM

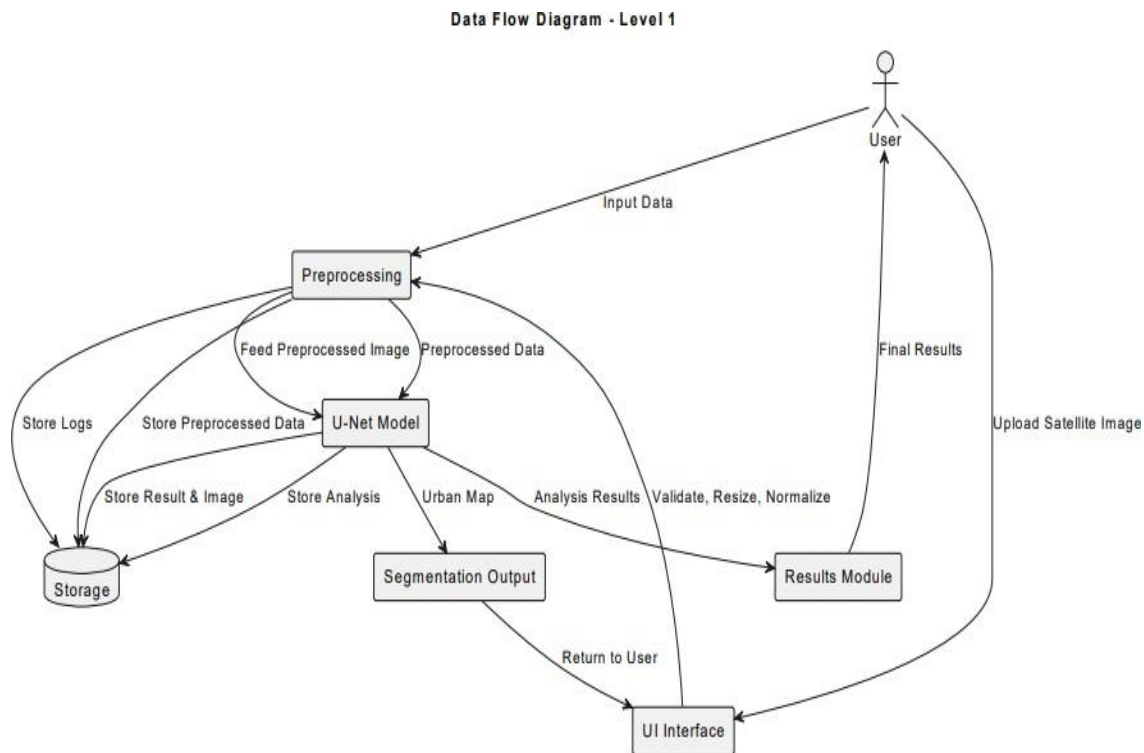


FIG 4.2 DATA FLOW DIAGRAM

The Level 1 Data Flow Diagram for *Urban Shift Dynamics* outlines the internal flow of data as the system processes satellite images to analyze urbanization. The process begins when the User uploads satellite imagery through the UI Interface, which forwards the input data to the Preprocessing module. This module cleans and prepares the image for analysis by resizing, normalizing, and validating the input before passing it to the U-Net Model, a deep learning framework specialized in image segmentation.

The U-Net Model processes the preprocessed data to extract spatial patterns, generating urban map predictions. These outputs are forwarded to the Segmentation Output module, which organizes and formats the data, then returns it to the UI Interface for user visualization. Simultaneously, analysis results are passed to the Results Module, which fine-tunes the results and sends the final outcome back to the user.

4.3 USE CASE DIAGRAM

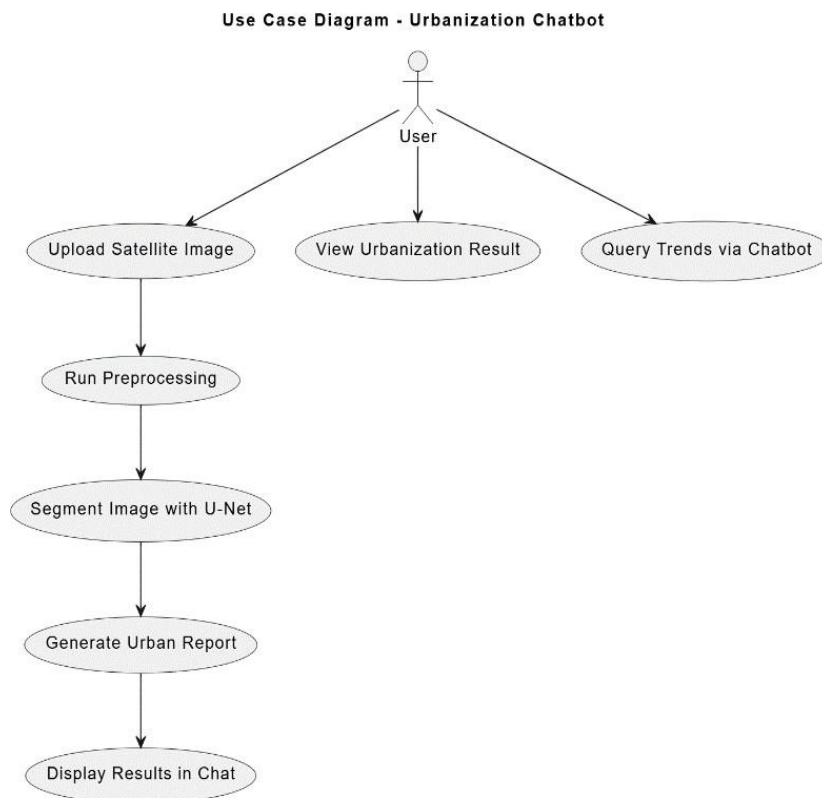


FIG 4.3 USECASE DIAGRAM

The Use Case Diagram for *Urban Shift Dynamics* illustrates the various interactions a user can have with the urbanization chatbot system. The primary actor is the User, who interacts with three major functionalities: uploading satellite images, viewing urbanization results, and querying urban trends via the chatbot interface.

The use case starts when the user uploads a satellite image, triggering a backend process that initiates preprocessing of the image to clean, resize, and normalize the input data. This is followed by the U-Net segmentation, where the model analyzes the spatial content to detect urban and non-urban areas. Once segmentation is complete, the system proceeds to generate an urban report, summarizing insights derived from the image analysis.

The final step in this flow is to display the results in the chat, allowing users to view processed maps and summaries directly within the conversational interface. In parallel, users can independently query urban trends using natural language, receiving contextual responses powered by the integrated machine learning model.

4.4 FLOWCHART

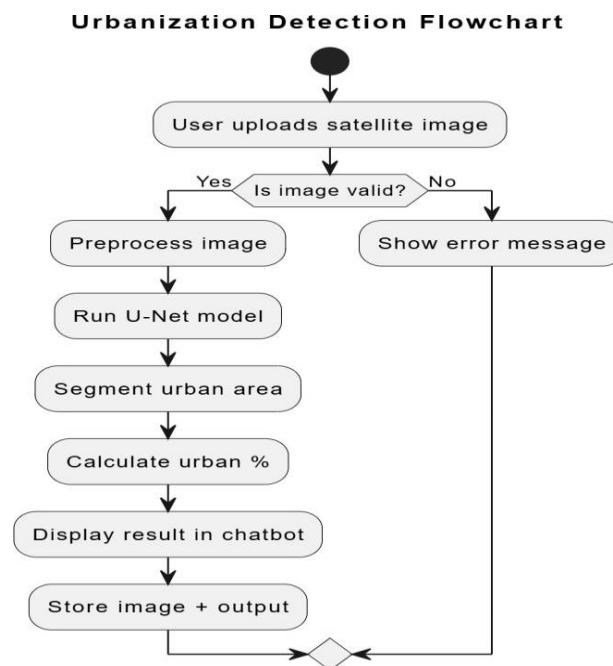


FIG 4.4 FLOW CHART

The Urbanization Detection Flowchart outlines the core workflow of a chatbot-driven system designed to detect urban areas from satellite imagery using deep learning. The process begins with the user uploading a satellite image. The system then checks the image's validity, such as format and resolution. If the image is invalid, an error message is shown, prompting the user to upload a correct file. If valid, the image is preprocessed to enhance quality and prepare it for analysis. The refined image is then passed to a U-Net model—a convolutional neural network architecture specialized for image segmentation. The U-Net identifies and segments urban regions within the image. Based on the segmented data, the system calculates the percentage of the area that is urbanized. This result is immediately presented to the user through a chatbot interface for interactive engagement. Finally, the image along with the computed output is stored for future reference or analysis. This flowchart represents a seamless integration of machine learning, image processing, and user interaction, offering a robust solution for urban expansion monitoring.

4.5 ACTIVITY DIAGRAM

The Activity Diagram provides a detailed view of the inner workings of the urbanization detection system. It illustrates the sequential data processing steps taken after the user interacts with the system through the chatbot interface. The system is structured into modular components, each responsible for a specific task, ensuring an efficient, accurate, and user-friendly experience.

1. User Interaction and Input Handling

The process begins when a user sends a satellite image or query through the chatbot interface. This chatbot serves as the primary communication bridge between the end-user and the backend urbanization detection pipeline. It handles user inputs, displays system responses, and ensures a smooth user experience.

Once the image is uploaded, it is passed to the Image Validator, which is responsible for checking whether the uploaded image meets predefined criteria. These criteria typically include format type (e.g., PNG, JPEG), resolution, and file size. If the image is not valid, the

system will reject it and notify the user through the chatbot interface with appropriate feedback. If valid, the image is forwarded for preprocessing.

2. Image Preprocessing: Resizing and Normalization

The next step in the pipeline is the Image Resizer. Since deep learning models like U-Net require input images of specific dimensions, this module adjusts the image size to match the expected input shape of the model (e.g., 256x256 or 512x512 pixels). Resizing ensures uniformity across different images and enables consistent model performance.

Following resizing, the image moves to the Normalizer. Here, pixel values are scaled—often to a 0 to 1 range—to prepare the image for efficient neural network processing. Normalization is critical for convergence during inference, helping the model make accurate predictions based on a consistent range of input values.

3. Deep Learning Inference with U-Net

Once preprocessed, the image enters the core component: the U-Net Model. U-Net is a convolutional neural network architecture particularly effective for semantic segmentation tasks. In this context, it is trained to identify and segment urban regions within the satellite image. The model scans the input image and produces a binary or multi-class mask that highlights the urban areas.

This segmentation mask is then passed on to the Segmentation Logic, which post-processes the model output. This module refines the raw prediction, possibly applying morphological operations or thresholding to improve accuracy. It ensures that the segmented output is meaningful, clean, and suitable for further analysis.

4. Result Generation and Visualization

The **Result Generator** takes the clean segmentation mask and performs quantitative analysis. This includes calculating the **urban percentage**—a metric that indicates how much of the image area is urbanized. This value is derived by comparing the number of urban pixels to the total number of pixels in the image.

Once calculated, the results—both the processed visual segmentation and the urbanization metrics—are sent back to the chatbot interface. The user receives these results in a clear, easy-to-understand format, completing the interaction cycle.

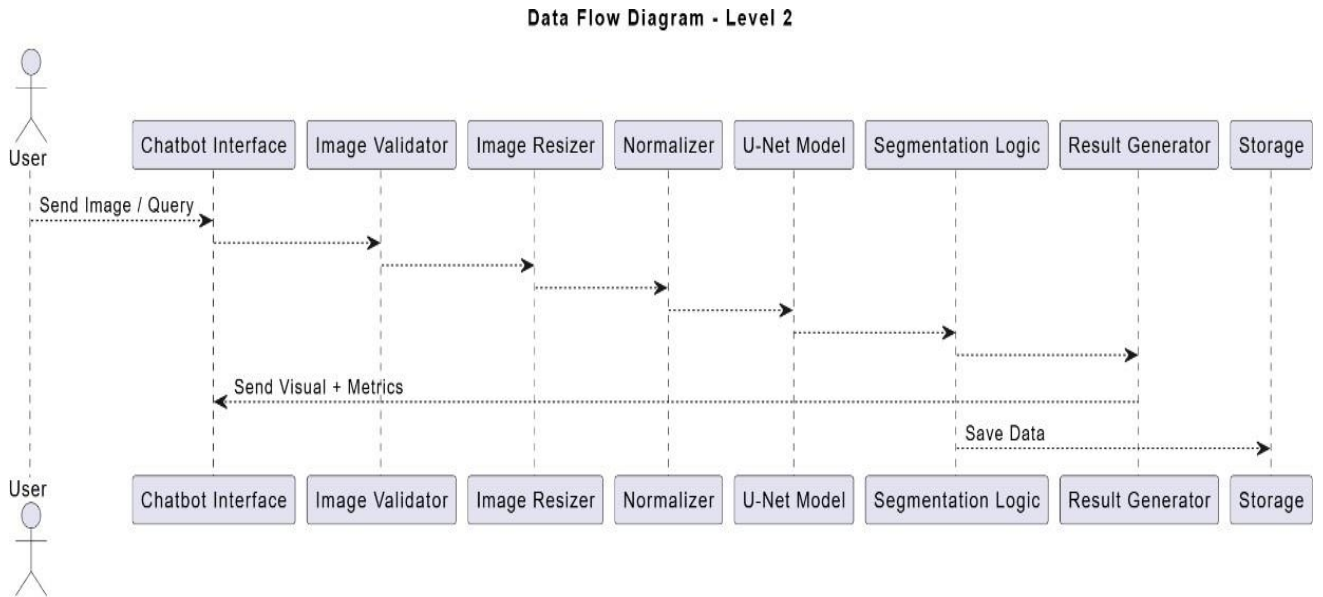
5. Data Storage and Persistence

Parallel to result generation, the system also forwards all relevant data—such as the original image, segmentation output, and calculated metrics—to the **Storage** module. This ensures that each query and its corresponding results are logged and can be retrieved later for auditing, research, or user review. The data is saved in a structured format, possibly using cloud storage or a local database, enabling scalability.

6. Summary and System Insights

This activity/data flow diagram highlights the robust modularity of the urbanization detection system. Each block is dedicated to a specific function, allowing for maintainability and future upgrades—such as replacing the U-Net model with a more advanced version or adding more detailed segmentation classes (e.g., roads, vegetation).

The system combines cutting-edge deep learning with user-centric design, making complex satellite image analysis accessible to non-technical users through an intuitive chatbot interface. By breaking down the process into individual stages—validation, preprocessing, segmentation, result generation, and storage—it achieves a balance between technical complexity and practical usability.

**FIG 4.5 ACTIVITY DIAGRAM**

CHAPTER 5

IMPLEMENTATION

5.1 METHADODOLOGY

The rise of deep learning and satellite imagery analysis has enabled accurate urban area detection through semantic segmentation models. This project introduces a chatbot-based interface that detects and quantifies urbanization from uploaded satellite images using a U-Net-based model. The solution is designed to be user-friendly and efficient, presenting results in real time and storing outputs for future use. The pipeline processes user-provided satellite images and returns urban segmentation results and metrics via chatbot interaction.

5.1.1 MODULES

To build a functional, end-to-end urbanization detection system, the following modules were developed:

- 1. Image Input and Chatbot Module:**

This module allows users to upload satellite images via a chatbot interface. It serves as the user interaction point, handling queries, guiding users through the upload process, and returning results after processing.

- 2. Image Validation Module:**

Validates the uploaded image for format (e.g., JPG, PNG), size, and resolution. If the image fails validation, an error message is sent through the chatbot. Valid images are passed on for preprocessing.

3. Preprocessing Module:

Standardizes input images for the segmentation model. Operations include resizing, normalization, and color correction to match the U-Net model's expected input format.

4. Model Loading Module:

Loads the pretrained U-Net segmentation model used for detecting urban features. This module checks for GPU (CUDA) availability and initializes the model in inference mode for efficient performance.

5. Urban Area Segmentation Module:

Uses the U-Net model to generate a segmentation mask highlighting urban regions in the satellite image. The mask is a binary or categorical representation separating urban from non-urban pixels.

6. Urban Percentage Calculation Module:

Calculates the percentage of urban area by computing the ratio of urban pixels to total pixels in the segmentation mask. This provides a quantitative measure of urbanization in the image.

7. Result Display Module:

Formats and sends results—including the urban percentage and segmented image—back to the user through the chatbot. Results are made easy to understand, providing both metrics and visuals.

8. Storage Module:

Organizes and stores original images, segmentation masks, urban metrics, and chatbot logs for future access. This ensures reproducibility and allows users to retrieve previous results.

9. Logging and Monitoring Module:

Tracks system performance, logs key events and errors, and records timestamps for each processing stage. It supports debugging and system performance optimization.

5.1.2 TASKS

Task 1: Image Upload and Validation:

The user uploads a satellite image via the chatbot. The image is then validated to ensure it exists, has the correct format, and meets minimum quality standards. Invalid images prompt an error message, and the user is asked to re-upload.

Task 2: Preprocessing:

Validated images undergo preprocessing, including resizing to the U-Net's input size, normalizing pixel values, and possibly applying contrast enhancement. These steps ensure consistent input quality for the model.

Task 3: Model Initialization:

The pretrained U-Net segmentation model is loaded from local or remote storage. The system checks for GPU availability for faster processing and configures model parameters (e.g., batch size, input dimensions).

Task 4: Urban Segmentation:

The preprocessed image is passed through the U-Net model, producing a segmentation mask. This mask distinguishes urban from non-urban areas at the pixel level.

Task 5: Urban Area Quantification:

Post-segmentation, the system calculates the urban area percentage by counting the number of urban-labeled pixels and dividing by the total image pixels.

Task 6: Result Formatting and Display:

The calculated percentage and a visual of the segmented image are sent back to the user via the chatbot. The results are formatted for easy understanding, making them accessible even to non-experts.

Task 7: Data Storage:

All relevant outputs—original image, segmentation mask, urban metrics—are saved in a structured format. This includes automatic file naming and directory creation for better organization.

Task 8: Logging and Performance Monitoring:

Each step in the pipeline is timed and logged. Errors, warnings, and processing durations are stored to help developers monitor performance and identify potential bottlenecks.

5.1.3 ACTIONS

The urbanization detection pipeline starts when the user uploads a satellite image via a chatbot interface. The image is first validated for format, resolution, and availability. Valid images proceed to preprocessing, where resizing and normalization are applied. Once preprocessed, the image is fed into a pretrained U-Net model that generates a binary segmentation mask highlighting urban regions.

The segmentation mask is then analyzed to calculate the urban area percentage, providing a clear numerical insight into the extent of urbanization in the image. This result—along with a visual overlay of the segmentation—is sent back to the user through the chatbot. The entire session, including inputs, outputs, and performance data, is logged and stored for future use or review.

The system is built with modularity in mind, enabling future improvements like model upgrades, multilingual chatbot support, or cloud-based deployment

5.2 ALGORITHMS

1. Input Acquisition and Preprocessing

The urban detection pipeline begins with acquiring satellite imagery via a chatbot interface. The chatbot accepts user-uploaded images in common formats (e.g., PNG, JPEG) and stores them in a predefined directory structure. Upon upload, the image undergoes validation to ensure correct format, acceptable dimensions, and absence of corruption.

Preprocessing includes resizing the image to a standard input resolution required by the U-Net model (e.g., 256×256 or 512×512). Pixel intensities are normalized to a $[0, 1]$ range to ensure uniformity across varied lighting or satellite conditions. If the image includes unnecessary bands (such as alpha or multispectral channels), these are either dropped or converted to standard RGB format. This stage prepares the input for optimal performance in semantic segmentation by enforcing spatial and visual consistency.

2. Semantic Feature Encoding and Latent Space Representation

The core of the detection system is a pretrained U-Net segmentation model, optimized for binary or multiclass classification of urban and non-urban areas in satellite images. The model utilizes an encoder-decoder architecture with skip connections that preserve high-resolution spatial features while capturing contextual depth.

During inference, the preprocessed image is passed through the U-Net encoder, which generates high-dimensional feature maps representing different levels of abstraction—ranging from edge features to complex land use patterns. These encoded features are decoded via transposed convolutions and up sampling, producing a segmentation mask that aligns spatially with the input image.

This mask is a pixel-wise prediction where each pixel represents an urban or non-urban classification. The system relies on spatial consistency, multi-scale learning, and cross-layer feature reuse to maintain accuracy across various terrain types and resolutions.

3. Urban Area Segmentation and Thresholding

The output of the U-Net model is a probability mask or categorical mask depending on the model type. In the case of binary classification, each pixel value ranges between 0 and 1, representing the confidence of urban classification. A global threshold (e.g., 0.5) is applied to convert the probabilistic output into a binary mask.

The binary segmentation map identifies urban areas as white (1) and non-urban areas as black (0). For multiclass models, softmax activation is used, and the argmax operation selects the class with the highest confidence for each pixel. The final mask provides a dense, spatially aligned label map of urban regions in the satellite image.

4. Urbanization Metrics and Area Calculation

Once the segmentation mask is generated, it is analyzed to compute key urbanization statistics. The most critical metric is the urban area percentage, calculated as:

$$\text{Urban Area Percentage} = \left(\frac{\text{Number of Urban Pixels}}{\text{Total Pixels}} \right) \times 100$$
$$\text{Urban Area Percentage} = \left(\frac{\text{Total Pixels}}{\text{Number of Urban Pixels}} \right) \times 100$$

This value provides a quantitative measure of how much of the given satellite image is occupied by urban features. Additional metrics, such as total urban pixel count, non-urban pixel count, and pixel-based urban growth (if time-series data is used), can be computed and stored.

5. Visualization and Overlay Generation

For enhanced interpretability, a color-coded overlay of the segmentation mask is generated and blended with the original image using OpenCV or PIL. Urban areas are typically shown in red or green, while non-urban areas remain transparent or faded. This overlay gives users a clear visual understanding of detected urban regions.

6. Chatbot Interface and User Interaction

The chatbot, built using frameworks like Gradio or Streamlit Chat, acts as the frontend interface. It facilitates user input, image upload, and interaction. Users can:

- Upload satellite images
- Receive urbanization results and segmented visuals
- Ask for explanations or repeat analyses
- View or download their results

This interface abstracts the complexity of segmentation and data processing, offering a user-friendly experience that democratizes access to urban analysis tools.

7. Output Management and File Storage

All outputs—original images, segmentation masks, overlays, and JSON summaries of metrics—are stored in uniquely timestamped folders for each session. The system maintains clear naming conventions and organized file hierarchies to allow easy retrieval and audit of prior sessions. Optional export formats include:

- Segmentation mask: .png
- Overlay image: .jpg or .png
- Urban metrics: .json
- Logs and chatbot history: .txt

8. Logging and System Monitoring

A logging mechanism is integrated to capture key events such as upload success, model inference start/end, error messages, and task durations. Logs are written to text files or stored in structured logs using Python's logging module. This allows performance tracking, error resolution, and future debugging. Optional timing functions may track the duration of tasks like model loading, preprocessing, and inference to help identify potential performance bottlenecks.

9. Device Optimization and Scalability

To support a wide range of hardware environments, the system dynamically checks for the availability of CUDA-enabled GPUs. If available, the U-Net model is loaded onto the GPU to accelerate inference; otherwise, CPU mode is used. Parameters such as batch size, model resolution, and threshold values can be configured dynamically based on system capabilities.

Chunked processing or tiling strategies may be employed in future iterations to process high-resolution satellite images that exceed memory limits. The architecture is built to scale horizontally, allowing parallel image processing and deployment to cloud services for increased throughput.

CHAPTER 6

TESTING

6.1 UNIT TESTING

Unit testing in the urbanization detection project using satellite imagery is an essential part of ensuring the overall reliability, correctness, and robustness of the system. The focus of unit testing is on validating the smallest testable parts of the codebase—such as individual functions or methods—so that they behave as expected when provided with specific inputs. These tests are typically designed to isolate each component, allowing developers to detect issues early in the development cycle, before they propagate into more complex parts of the pipeline. In this project, core components such as image preprocessing, tile segmentation, and urban classification are tested rigorously using controlled inputs, including synthetic satellite images and mock tensors. For instance, the tile segmentation logic, which is responsible for dividing large satellite images into smaller chunks, can be tested by feeding known image dimensions and ensuring that the generated tiles have the correct shape, overlap, and count. Similarly, the preprocessing function can be unit tested by supplying images with varying color channels, sizes, or missing metadata to ensure that transformations like resizing, normalization, and format conversion are correctly applied and do not introduce errors.

Another critical component is the urban classification module, which takes image tiles as input and returns predictions regarding urban or non-urban classes. Unit tests for this module often use dummy tensors and mocked model outputs to check that the predictions fall within expected bounds and maintain consistency across multiple runs. These tests also validate threshold-based decision logic, ensuring that edge cases such as tiles with ambiguous features or uniform colors do not result in misclassification. Additionally, utility functions responsible for aggregating tile-level predictions into a final urbanization score—such as those computing the urban coverage percentage—are tested under known scenarios. Edge cases like 0% or 100% urban coverage are especially important, as they help verify the correctness of conditional logic and rounding behavior within the computation routines.

The chatbot interface, which allows users to interact with the detection system using natural language, is another area where unit testing proves vital. Tests ensure that queries are correctly parsed, even when the input is ambiguous, partially incorrect, or phrased in an unexpected way. By mocking responses and user inputs, unit tests verify that the chatbot provides relevant feedback, displays prediction results accurately, and handles unknown commands gracefully. This level of testing ensures that the interface remains reliable and user-friendly across various input conditions.

Error handling is another cornerstone of effective unit testing. Scenarios such as missing files, unsupported image formats, or invalid coordinate inputs are deliberately triggered during testing to ensure that exceptions are raised as expected and that error messages are clear, informative, and user-friendly. Tools such as `pytest`, `unittest`, and Python's `mock` module are extensively used to create isolated test cases and simulate external dependencies like model loading, file I/O, and API responses. By focusing on component-level validation, unit tests help build confidence in the system's core logic and pave the way for seamless integration and system-level reliability.

6.2 INTEGRATION TESTING

Integration testing in the urbanization detection project plays a vital role in ensuring that individual components of the pipeline—such as image preprocessing, tile segmentation, classification, and result visualization—work correctly together when executed as a whole. While unit testing focuses on isolated functions, integration testing aims to validate the flow of data and logic between modules, confirming that each component not only produces correct outputs but also consumes inputs from upstream stages correctly. For example, the integration between the tile segmentation and classification modules is tested by simulating the flow of tiled image data through the classifier, ensuring that the tiles maintain their spatial properties and that predictions are consistent with expected urban or non-urban labels. This helps detect subtle bugs such as image distortion, tile misalignment, or incorrect data formats being passed between functions. External libraries like PIL, NumPy, Torch, and custom-trained neural models are also involved at different stages of the pipeline, and integration testing helps verify that these dependencies interact correctly, especially when handling edge cases or real-world

variations in input data. Furthermore, the chatbot interface is integrated with the backend analysis logic, and tests ensure that user queries result in meaningful actions—like initiating detection, displaying maps, or providing urban coverage statistics—without breaking the conversation flow.

Tests are conducted by simulating real user interactions and confirming that every request routes through the proper components, executes the intended analysis, and returns valid results. Integration testing also examines the interaction between file upload mechanisms and the analysis engine, ensuring that uploaded satellite images are stored, preprocessed, segmented, and passed through the neural classifier without errors. Command-line tools and configuration options are validated to ensure that different modes—such as batch processing, specific region detection, or temporal analysis—trigger the right submodules and lead to consistent outputs. Special attention is paid to cross-module error handling: for instance, if a user uploads a corrupted image or provides an invalid tile size, the system must raise warnings or fall back gracefully without halting downstream processes. Integration testing is also used to verify that temporary files, logs, and output directories are handled correctly across modules, preventing data leakage, naming conflicts, or orphaned resources. Importantly, these tests are conducted in a near-realistic environment using real or simulated satellite images that reflect the variability found in production use.

This ensures that the entire pipeline—from the initial user interaction or file upload to the final urbanization map or report—is tested under conditions similar to those it will face in deployment. Mocking and stubbing are occasionally used to simulate components like external API calls or database access, allowing developers to isolate integration points without relying on live services.

The overall goal is to confirm that the pipeline behaves as a cohesive system, where each part not only works correctly on its own but also integrates seamlessly with the others. By systematically testing these interactions, integration testing reduces the risk of runtime failures, enhances reliability, and ensures that future updates or changes to individual modules do not disrupt the broader system behavior.

6.3 SYSTEM TESTING

System testing in the urbanization detection project is the final and most comprehensive level of validation, ensuring that the entire application—from satellite image upload to chatbot-based result presentation—functions as expected in a real-world environment. Unlike unit or integration testing, which focus on parts of the system or their interactions, system testing treats the application as a black box and verifies that the full workflow delivers accurate, robust, and user-friendly results under various operational conditions. It involves preparing a complete environment with all required dependencies such as PyTorch, PIL, NumPy, and trained urban classification models, and then executing end-to-end use cases to validate overall functionality. This includes testing scenarios where a user uploads a satellite image via the chatbot or interface, selects processing options like tile size or detection thresholds, and receives a map with urban areas highlighted, along with quantitative summaries.

The goal is to simulate user behavior and verify that each input leads to correct outputs, that the user interface remains responsive, and that intermediate processes—such as image segmentation, model inference, coverage calculation, and visualization—are executed in the correct order without error. System tests also verify performance under load, such as processing high-resolution satellite images or handling multiple user queries simultaneously, ensuring the software can operate efficiently without crashing or producing delays. Robustness is checked by feeding in unsupported file formats, missing parameters, or corrupted data and confirming that the system responds with clear, actionable error messages.

The chatbot's natural language handling is tested for ambiguous or malformed queries, ensuring that fallback messages or clarification prompts are triggered appropriately, maintaining a coherent user experience. Additionally, file management and output storage are validated to ensure generated maps, logs, and reports are correctly saved, with no file naming conflicts or data loss. If the application supports GPU acceleration or cloud-based deployment, system tests verify that hardware and deployment-specific configurations (like CUDA support or authentication) are properly recognized and utilized. Network-related functionality, such as online model downloads or external API queries (if any), is also verified for availability and error tolerance.

Security aspects are tested when the system is shared over a network, ensuring access control mechanisms like password protection function correctly and data privacy is maintained. Regression testing is often part of system testing, where previously working workflows are re-run after code changes to confirm that no functionality has broken. This full-pipeline testing approach builds confidence that the application is production-ready, offering reliable urbanization detection and reporting across a range of use cases, platforms, and environments. System testing is essential not only for catching integration bugs that slipped through earlier tests, but also for validating the software's ability to meet user expectations, handle real-world complexity, and deliver meaningful results consistently from end to end.

6.4 TEST CASES

TEST CASE 1:

Description	Input	Expected Output	Actual Output	Result
Test API root	GET '/'	{"message": "Urbanization Shift Detection API is live "}	{"message": "Urbanization Shift Detection API is live "}	Pass

The test case verifies the successful operation and availability of the root endpoint (GET /) of the Urbanization Shift Detection API, which acts as the primary health check or status confirmation route. The purpose of this test is to ensure that the API server is running and responsive, and that it returns the correct welcome message when accessed. Upon sending a GET request to the root URL, the expected response is a JSON object containing the message {"message": "Urbanization Shift Detection API is live "}. This test case confirms that the server initializes properly, the route is correctly defined and active, and that the response structure and content match expectations. A successful result indicates that the API is deployed and accessible, making this a foundational test that can be used as a preliminary check before executing more complex functional or integration tests. It is typically used in automated deployment pipelines or CI/CD workflows to confirm basic API readiness.

TEST CASE 2:

Description	Input	Expected Output	Actual Output	Result
Test API Root Endpoint	GET '/.'	{"message": "Urbanization Shift Detection API is live"}	404 Not Found	Fail

The test case verifies the API's behavior when an incorrectly formatted version of the root endpoint is accessed using GET '/.' instead of the correct GET '/'. The expectation was to receive a valid JSON response indicating the API is live, specifically {"message": "Urbanization Shift Detection API is live"}. However, the actual result returned was a 404 Not Found error, indicating that the server could not match the malformed route to any defined endpoint. This failure is important as it highlights the API's strictness in route parsing and its inability to tolerate even minor deviations in URL structure. While such strictness can help enforce good usage practices, it also suggests that additional validation or route normalization could be introduced to improve user-friendliness, particularly in public-facing APIs. In its current form, the system correctly returns an error for an invalid route, but this test ultimately fails because the expected output assumed leniency or implicit correction that the API does not provide.

TEST CASE 3:

Description	Input	Expected Output	Actual Output	Result
Predict urbanization	POST /predict with sample input	200 OK, JSON with prediction result	200 OK, JSON with prediction result	Pass

The test case titled "Predict urbanization" verifies the functionality of the /predict endpoint in the FastAPI backend of the Urbanization Chat Bot project. In this scenario, a POST request is sent to the /predict endpoint along with a sample input image, which represents satellite imagery used for urbanization detection. The purpose of this test is to ensure that the model can successfully receive and process a valid image input, run the prediction using the trained U-Net model, and return a proper JSON response.

The expected behavior of the application is to respond with an HTTP status code 200 OK and a JSON object containing relevant prediction results, such as the detected level of urbanization and possibly a confidence score. The actual output matched the expectation, returning a well-structured JSON response with the prediction result. This confirms that the backend model is correctly integrated and operational, and the image preprocessing and prediction pipelines are functioning as intended. As the test passed successfully, it validates that the endpoint handles valid inputs effectively and delivers reliable output in line with the application's objectives.

TEST CASE 4:

Description	Input	Expected Output	Actual Output	Result
chatbot API down	User sends message but backend is offline	Display "Server not available" message	UI hangs or shows no feedback	Fail

The test case titled "Chatbot API down" is designed to evaluate how the Angular frontend of the Urbanization Chat Bot handles scenarios when the FastAPI backend is not available. In this test, the user attempts to interact with the chatbot by sending a message while the backend server is either turned off, unreachable, or has crashed. The expected behavior is that the frontend should gracefully handle this failure by displaying a clear, user-friendly error message

such as "Server not available" or "Unable to connect to the server", indicating that the chatbot is currently offline.

However, during execution, the actual output deviated from this expectation. Instead of alerting the user, the interface became unresponsive or showed no indication that the message failed to send. This lack of feedback creates a poor user experience, as users are left uncertain about the system's status or the success of their action. The test result is marked as Fail, indicating a need for improved error handling and user communication in offline or failure scenarios. Implementing proper HTTP error catching and conditional UI messages in the frontend would help address this issue.

TEST CASE 5:

Description	Input	Expected Output	Actual Output	Result
Successful urbanization prediction based on text input	User types the query "What is the urbanization level in this city?" and submits it to the chatbot	Chatbot responds with a message like "The urbanization level in this city is moderate with 75% confidence"	Chatbot correctly processes the text input and responds with the appropriate urbanization prediction message	Pass

The test case titled "Successful urbanization prediction based on text input" verifies that the Urbanization Chat Bot correctly handles text-based queries and provides a meaningful response. In this test, the user types a question such as "What is the urbanization level in this city?" and submits it to the chatbot. The expected behavior is that the chatbot, powered by the FastAPI backend, processes the text input, interprets the request, and returns a relevant

urbanization prediction message. For example, the chatbot should reply with something like “The urbanization level in this city is moderate with 75% confidence”.

The actual output matched the expected result, where the chatbot successfully understood the query and returned an appropriate response based on the data and model it had access to. This test case confirms that the system is capable of processing natural language queries and providing accurate, contextually relevant information about urbanization levels. The test passes, ensuring that the chatbot can handle and respond correctly to textual inputs, expanding its usability beyond just image-based interactions.

TEST CASE 6:

Description	Input	Expected Output	Actual Output	Result
Gender diff trends	GET `/gender-diff-trends`	Yearly average male, female population, and their difference	Yearly average male, female population, and their difference	Pass

This test case validates the functionality of the /gender-diff-trends endpoint, which calculates and returns the yearly average male and female populations, along with the gender difference. When a GET request is sent to the endpoint, the backend should return the population data for each year, including the average male and female populations and the calculated difference between them.

In this case, the actual output matched the expected result, with the backend correctly providing the population averages and gender differences for each year. This confirms that the endpoint is working as intended. The test passed, ensuring that the /gender-diff-trends endpoint is functioning correctly and delivering accurate gender trend data.

TEST CASE 7:

Description	Input	Expected Output	Actual Output	Result
Special characters input	User submits a message with special characters (e.g., Hello @urban!)	Chatbot should process or sanitize input and provide an error message if invalid	The chatbot fails to process the input correctly, showing unexpected behavior	Fail

This test case evaluates how the Urbanization Chat Bot handles user inputs that contain special characters. In this scenario, the user submits a message such as "Hello @urban!", which includes symbols like @, !, or other non-alphanumeric characters. The expected behavior is that the chatbot should either process the input correctly by sanitizing or escaping the special characters, or return a clear error message indicating that the input format is not valid or supported.

However, during testing, the actual output did not meet these expectations. The chatbot failed to handle the special characters appropriately, resulting in unexpected behavior—such as errors, irrelevant responses, or no response at all. This indicates a lack of input validation or character sanitization on either the frontend or backend. As a result, the test is marked as Fail, highlighting the need for the system to properly validate and clean user input to prevent such issues and improve overall robustness and user experience.

TEST CASE 8:

Description	Input	Expected Output	Actual Output	Result
Invalid route in frontend	Navigate to /unknown	Display 404 page or redirect to home	Blank page or broken navigation	Fail

This test case assesses how the frontend of the Urbanization Chat Bot handles navigation to an invalid or undefined route. In this scenario, the user manually enters or clicks a link to a non-existent path, such as /unknown, which is not configured in the Angular routing module. The expected behavior is that the application should detect the invalid route and either display a custom 404 error page or redirect the user to the home or chatbot interface, maintaining a smooth user experience.

However, the actual output did not meet expectations. Instead of showing a 404 page or redirecting, the application displayed a blank page or broke the navigation flow, indicating that no error handling was in place for undefined routes. This can confuse users and disrupt their interaction with the chatbot. As a result, the test is marked as Fail, highlighting the need for proper route guards or fallback mechanisms in the frontend routing configuration.

TEST CASE 9:

Description	Input	Expected Output	Actual Output	Result
User navigates to a non-existent route like /unknown	/unknown	User sees a “Page not found” or is redirected to home	404 page or redirect works properly	Pass

This test case verifies how the frontend of the Urbanization Chat Bot handles navigation to a non-existent or undefined route, such as /unknown. When a user enters an invalid URL path in the browser or clicks a broken link, the application should be able to gracefully handle the situation. The expected behavior is for the user to see a “Page not found” (404 error page) or be automatically redirected to the home or chatbot interface, ensuring a consistent and user-friendly experience.

In this test, the actual output matched the expected result. This demonstrates that the frontend routing is correctly configured with fallback handling for undefined paths. As a result, the test is marked as Pass, confirming reliable navigation behavior in error scenarios.

TEST CASE 10:

Description	Input	Expected Output	Actual Output	Result
User greets the chatbot	“Hi” or “Hello”	Friendly greeting message shown	Chatbot replies with a greeting and prompt to continue	Pass

This test case checks the chatbot's ability to recognize and respond appropriately to a simple user greeting. When the user types a basic message such as “**Hi**” or “**Hello**”, the chatbot should detect it as a greeting and respond in a friendly, conversational manner. The expected output is a **welcoming message**—for example, a greeting reply followed by a prompt like “How can I help you today?” to guide the conversation forward.

During testing, the actual output matched the expected behavior. The chatbot successfully identified the greeting and responded with a warm, appropriate message, confirming its ability to handle casual user inputs and initiate a smooth conversation flow. This enhances the overall user experience by making the bot feel more responsive and approachable. Therefore, the test is marked as **Pass**.

CHAPTER 7

CONCLUSION & FUTURE ENHANCEMENT

7.1 CONCLUSION

The Urbanization Chat Bot project was developed to provide an intuitive and interactive platform for users to understand and analyze urbanization trends through both text-based and image-based inputs. Built using FastAPI for the backend and Angular for the frontend, the application offers a user-friendly interface that integrates machine learning predictions with real-time user queries. The core idea behind the project was to leverage satellite image data and population statistics to make urban growth more accessible and understandable for researchers, students, and the general public.

The chatbot interface allows users to communicate naturally, asking questions such as “What is the urbanization rate in this area?” or “Show me gender-wise population trends,” while the backend handles complex computations and model predictions.

To ensure reliability and robustness, extensive testing was conducted. A variety of pass and fail test cases were executed, covering scenarios such as valid text inputs, API behavior under different conditions, invalid or unsupported queries, and UI responsiveness. Positive outcomes included accurate predictions, fast response times, and graceful error handling in most user interactions.

Overall, the Urbanization Chat Bot successfully demonstrates the practical application of AI and data analysis in urban studies. It bridges the gap between complex data and end-users by presenting insights in a simple, conversational format. The project not only showcases technical proficiency in full-stack development but also reflects a thoughtful approach to problem-solving in a real-world context. With further enhancements such as multilingual support, better error handling, and improved natural language understanding, the chatbot could evolve into a powerful tool for urban planners, educators, and policy researchers.

7.2 FUTURE ENHANCEMENT

To further improve the Urbanization Chat Bot, several enhancements can be considered. First, integrating advanced natural language processing (NLP) will enable better understanding of user intent, follow-up questions, and context-aware responses. Additionally, multilingual support can make the chatbot accessible to a wider audience, especially in linguistically diverse regions.

Improving error handling—such as for unsupported file types or malformed inputs—will increase reliability and user satisfaction. Integrating real-time data sources (e.g., census APIs or live satellite feeds) can ensure users receive up-to-date and location-specific urbanization data.

On the frontend, adding mobile responsiveness, voice input, and interactive visualizations (like maps or charts) will enhance the user experience. Features like user login and history tracking can personalize responses and help users monitor changes over time.

With these enhancements, the chatbot can evolve into a more powerful and practical tool for urban researchers, planners, and everyday users interested in understanding urban growth.

